M-CORE — Prior Art & Method Disclosure (High-Level)

Status: Public Method Disclosure (Prior Art)
Version: 1.0
Date: 22 May 2024
Author: Yuri S. Timofeev
Associated Deposits: n'RIS / RCIS (IDs referenced separately)

---

1. Purpose and Scope

This document provides a **high-level public disclosure** of the methodological principles underlying **M-CORE**, an external cognitive control architecture designed to operate alongside large language models (LLMs).

The purpose of this disclosure is to:

- establish **authorship and development priority**,
- define the **conceptual scope** of the M-CORE approach,
- and document an independent line of research into externally governed, human-directed AI systems.

This disclosure is **non-exhaustive** and intentionally omits implementation-specific details.

---

2. Terminology and Definitions

- **External Evaluation Core (EEC):**
  A modular control layer operating outside the neural weights of an LLM, responsible for evaluation, filtering, and long-term behavioral regulation.
- **Behavioral Valuation Ledger (BVL):**
  A persistent system of evaluative signals (+/−) used to record outcomes of interactions and construct a long-term preference and constraint structure. This ledger replaces transient conversational history with consequence-based memory.
- **Human-Defined Objective Layer (HDOL):**
  An interface through which human-defined goals, constraints, and priorities are translated into operational evaluation criteria for the system.
- **Multi-Agent Collision Protocol (MAC):**
  A structured interaction framework in which multiple specialized AI instances operate toward a shared objective under expert human oversight, enabling cross-validation and reduction of systemic errors.

---

3. Core Methodological Claims

This disclosure asserts the following high-level methodological principles:

1. **Separation of Model and Governance:**
   The behavioral regulation, long-term objectives, and evaluative logic of an AI system can be maintained independently from the underlying language model.
2. **Consequence-Based Memory:**
   System behavior is shaped by persistent evaluative records of outcomes rather than by raw conversational transcripts.
3. **Domain-Specific External Cores:**
   Specialized professional behaviors can be instantiated by modifying the external evaluation core without retraining the underlying model.
4. **Structured Sensor Pre-Processing (Conceptual):**
   Non-textual inputs (e.g., visual or acoustic signals) may be transformed into structured symbolic representations prior to interaction with an LLM.

---

4. Non-Disclosure and Limitations

This document intentionally excludes:

- implementation details,
- data schemas,
- cryptographic methods,
- internal evaluation logic,
- and integration specifics.

Such elements constitute unpublished research and are not disclosed in this public statement.

---

5. Intellectual Property Notice

This disclosure does not grant any license or right to use, reproduce, or commercialize the described methodology.

This document serves solely as a public record of conceptual development and prior art.

This disclosure is provided for documentation and research transparency purposes.